

Create training job

When you create a training job, Amazon SageMaker sets up the distributed compute cluster, performs the training, and deletes the cluster when training has completed. The resulting model artifacts are stored in the location you specified when you created the training job. [Learn more](#)

Job settings

Job name

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

IAM role

Amazon SageMaker requires permissions to call other services on your behalf. Choose a role or let us create a role that has the [AmazonSageMakerFullAccess](#) IAM policy attached.

AmazonSageMakerUserIAMExecutionRole_ce175b72 ▾

Create role using the role creation wizard

Algorithm options

Use an Amazon SageMaker built-in algorithm, your own algorithm, or a third-party algorithm from AWS Marketplace.

Algorithm source

- ☒ Amazon SageMaker built-in algorithm [Learn more](#)
- ☐ Your own algorithm resource
- ☐ Your own algorithm container in ECR [Learn more](#)
- ☐ An algorithm subscription from AWS Marketplace

Choose an algorithm

Choose an algorithm or custom training image... ▾

☐ Enable SageMaker metrics time series

Resource configuration

Instance type

ml.m4.xlarge ▾

Additional storage volume per instance (GB)

1

Instance count

Specify the number of instances that you want to allocate to this training job.

1

Keep alive period

Use [SageMaker Training Managed Warm Pools](#)

minutes ▾

Maximum runtime is 1 hour (60 minutes or 3600 seconds).

Encryption key - optional

Encrypt your data. Choose an existing KMS key or enter a key's ARN.

No custom encryption ▾

Stopping condition

Maximum runtime

24

hours ▾

Network

aws

Q

[Alt+S]

Marketplace products.

VPC - optional

For better security, we recommend that you use a private VPC.

Select a VPC

Hyperparameters

Choose an algorithm or custom training image to get started.

Input data configuration

Create up to 20 channels of input sources. If the algorithm you chose supports multiple input channels, you can specify those here. See [Algorithms Provided by Amazon SageMaker: Common Parameters](#)

Channels

▼ train

Remove

Channel name

train

Input mode - optional

Content type - optional

Compression type

None

Record wrapper

None

Data source

☒ S3

☐ File system

S3 data type

S3Prefix

S3 data distribution type

FullyReplicated

S3 location

s3://bucket/path-to-your-data/

Add channel

Checkpoint configuration - optional

The algorithm is responsible for periodically generating checkpoints. The checkpoints are saved to this location and used to resume managed spot training jobs. [Learn more.](#)

S3 location

s3://bucket/path-to-your-data/

To find a path, [go to Amazon S3](#)

Local path - optional

If the algorithm provides checkpoints, this is the local location the checkpoints are written to. If you do not specify a location, the checkpoints are written to /opt/ml/checkpoints.

Output data configuration

S3 output path

s3://bucket/path-to-your-data/

Encryption key - optional

If you want Amazon SageMaker to encrypt the output of your training job using your own AWS KMS encryption key instead of the default S3 service key, provide its ID or ARN.

Privacy

Terms

© 2026, Amazon Web Services, Inc. or its affiliates.

https://ap-south-1.console.aws.amazon.com/sagemaker/home?region=ap-south-1#/jobs/create

2/3

aws

Q

[Alt+S]

Enable managed spot training - *optional*

Save compute costs for jobs that have flexibility in start and end times. Amazon SageMaker will use spare capacity only to run this job. [Learn more](#)

Maximum wait time before job terminates *optional stopping condition*

At the end of this duration you will receive the complete or partial results of you managed spot training job.

48

hours

▼ Tags - *optional*

Key	Value	
		Remove

Add tag

Cancel

Create training job